

Assignments 7, 8, and 9, due May 10th, 17th, and 24th 2019

Project “Count Sort”

Algorithm:

The count sort algorithm consists of three steps. Start with an array of unsorted integer numbers:

A[0]	A[1]	A[2]	A[3]	A[4]	A[5]	A[6]	A[7]
2	5	3	0	2	3	0	3

- Create a histogram of the occurrence of each number:

C[0]	C[1]	C[2]	C[3]	C[4]	C[5]
2	0	2	3	0	1

- For each element c in C , compute the number of elements in A which are less or equal to c (prefix sum shifted by 1):

C[0]	C[1]	C[2]	C[3]	C[4]	C[5]
0	2	2	4	7	7

- Copy each value a in A into a new array B . The target position in B is determined by the corresponding number for a in C . After each copy operation, the corresponding number in C will incremented, thus writing the next element of the same value in the adjacent position:

A[0]	A[1]	A[2]	A[3]	A[4]	A[5]	A[6]	A[7]	C[0]	C[1]	C[2]	C[3]	C[4]	C[5]	B[0]	B[1]	B[2]	B[3]	B[4]	B[5]	B[6]	B[7]
2	5	3	0	2	3	0	3	0	2	2	4	7	7			2					
2	5	3	0	2	3	0	3	0	2	3	4	7	7			2					5
2	5	3	0	2	3	0	3	0	2	3	4	7	8			2		3			5
2	5	3	0	2	3	0	3	0	2	3	5	7	8	0		2		3			5
2	5	3	0	2	3	0	3	1	2	3	5	7	8	0		2	2	3			5
2	5	3	0	2	3	0	3	1	2	4	5	7	8	0		2	2	3	3		5
2	5	3	0	2	3	0	3	1	2	4	6	7	8	0	0	2	2	3	3		5
2	5	3	0	2	3	0	3	2	2	4	6	7	8	0	0	2	2	3	3	3	5

Use case:

- For these assignments, we sort a list of (age | name) entries according to age. Age is represented as an `int` with a value range from 0 to 120, name is represented as `char[32]`.

First assignment (until May 10th)

Assignment:

- Write a sequential C program `list_gen.c` that creates and prints a list of N random (age|name) entries, with N being a command line parameter. The seed for the random number generator should be a second command line parameter. For generating random names, use the provided function `gen_name(const char[32] buffer)` in header file `people.h`.

Example:

```
[philipp@hudson sort]$ ./list_gen 5 1337
```

```
41 | Isela Stayer
27 | Lanell Benberry
61 | Chaya Kijek
35 | Malika Harbour
44 | Mitzi Mccalley
```

- Implement a sequential version of the count sort algorithm for sorting the generated list on the CPU. The program `list_sort.c` should take the same parameters as `list_gen.c` and print the list of names twice, first unsorted and then sorted.
- Come up with a parallelization strategy for every step of the count sort algorithm. Consider factors such as correctness, synchronization, and performance on GPUs and CPUs. Write down your parallelization strategy and considerations into a file `parallelization_plan.txt`.

Resources:

- `people.zip`
Holds a header file `people.h` and text files for generating names.

Hints:

- In order to support a structured approach, the data type for the (age|name) entries is defined in `people.h`:

```
typedef struct {
    int age;
    name_t name;
} person_t;
```

This data structure must be used for all exercises!
- Organize your program into appropriate functions, for example:

```
void generate_list(person_t** list, int entries) { ... }
void print_list(person_t* list, int entries) { ... }
```

Also, a structure resembling the individual steps of the algorithm is required for the second assignment.

Solution upload:

- `list_gen.c`, `list_sort.c`, `parallelization_plan.txt`
- Via e-mail to philipp.gschwandtner@uibk.ac.at – one submission per group only!
Subject: “[PS703106] [AS07] GR_## - NAME1, NAME2, NAME3”
Solution must be submitted before Friday May 10th, 09:15!

Second assignment (until May 17th)

Assignment:

- Implement an OpenCL version of the prefix sum algorithm (second step of the count sort algorithm) as discussed in the lecture:
 - Prefix sum implementation for a single work group according to Hillis and Steele (hillisstele.{c/cl}).
 - Optimized implementation using down-sweep step (downsweep.{c/cl}).
 - Extension to arbitrarily large arrays using multiple work groups (prefixglobal.{c/cl}).
- For each of the aforementioned variants, a test program must be provided.

Solution upload:

- hillistele.{c/cl}, downsweep.{c/cl}, prefixglobal.{c/cl}
- Via e-mail to philipp.gschwandtner@uibk.ac.at – one submission per group only!
Subject: “[PS703106] [AS08] GR_## - NAME1, NAME2, NAME3”
Solution must be submitted before Friday May 17th, 09:15!

Third assignment (until May 24th)

Assignment:

- Implement list_sort of the first assignment in OpenCL using the prefix sum implementation of the second assignment.
 - Implement the histogram step (first step of count sort) using OpenCL
 - Implement the element copying (third step of count sort) using OpenCL
 - Write a test program countsort_bench.c, which takes the same parameters as the program of the first assignment. However, instead of printing the names, it should just measure and print the time required for sorting.

Example:

```
[philipp@hudson sort]$ ./countsort_bench 2000000 1337  
Sequential: 12.11 ms  
OCL: 4.32 ms
```

Solution upload:

- countsort_bench.c, including all involved source/header files
- Via e-mail to philipp.gschwandtner@uibk.ac.at – one submission per group only!
Subject: “[PS703106] [AS09] GR_## - NAME1, NAME2, NAME3”
Solution must be submitted before Friday May 24th, 09:15!